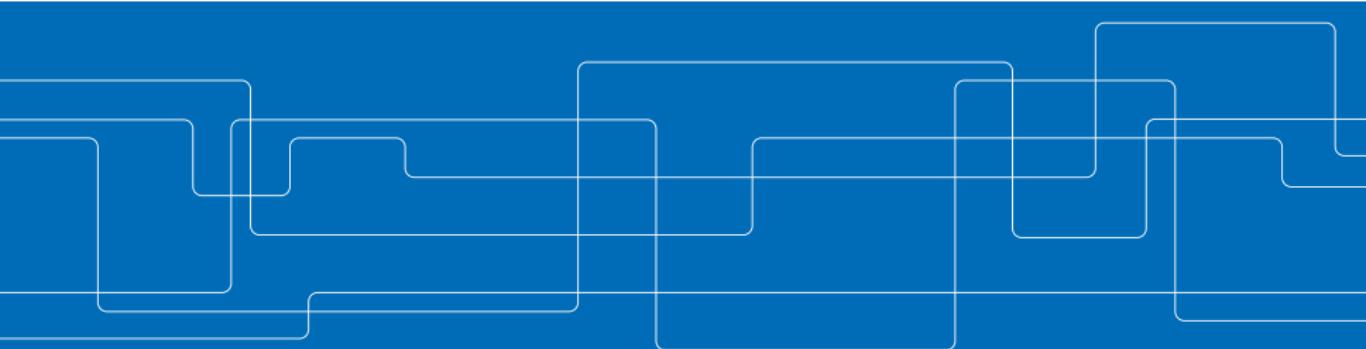




Designing Neural Network Architectures Using Reinforcement Learning

Authors : Bowen Baker, Otkrist Gupta, Nikhil Naik and Rameh Raskar

Presenter: Ezgi Korkmaz





CNN and RL

- ▶ A typical CNN consists of several convolution, pooling and fully connected layers.
- ▶ Numerous design choices for CNN architecture
 - ▶ The number of layers each type
 - ▶ The ordering of layers
 - ▶ The hyperparameters of the each type of layer.
- ▶ The design space of CNN architecture is extremely large
- ▶ In this paper automate the process of CNN architecture selection based on reinforcement learning
- ▶ Construct a Q-learning agent whose goal to discover CNN architecture which performs well
 - ▶ The agent learns through random exploration finds the higher performing models using the ϵ -greedy strategy.
 - ▶ Agents gets a validation accuracy score instead of a reward



CNN and RL

- ▶ Construct a Q-learning agent whose goal to discover CNN architecture which performs well
 - ▶ The agent learns through random exploration finds the higher performing models using the ϵ -greedy strategy.
 - ▶ Agents gets a validation accuracy score instead of a reward
 - ▶ Expedited learning process through repeated memory sampling using experience replay.
 - ▶ Q-learning based meta-modelling method as MetaQNN



Desingning Neural Network Architectures

- ▶ Related work
 - ▶ 1980s genetic algorithm based approaches [Schaffer et al. 1992]
 - ▶ Bayesian optimization for network architecture and [Bergstra et al. 2013]
 - ▶ Bayesian optimization for hyperparameters [Shahriari et al. 2016]
 - ▶ Meta modelling approach based on Tree of Parzen Estimators (TPE) [Bergstra et al. 2011]
- ▶ All failed to match the performance of handcrafted networks.



Q-Learning Background

- ▶ The task of teaching an agent to find an optimal paths as a Markov Decision Process (MDP) in a finite horizon environment.
- ▶ Environment will have a discrete finite state and action space S and U
- ▶ For any state $s_i \in S$ there is a finite set of actions $U(s_i) \subseteq U$
- ▶ Transitions are stochastic transition from s_i to s_j by taking $u \in U(s_i)$ is $p_{s'|s,u}(s_j|s_i, u)$
- ▶ At each time step t the reward r_t can be stochastic
- ▶ Agents goal to maximize the total expected reward over all possible trajectories $\max_{\tau_i \in \tau} R_{\tau_i}$



Q-Learning Background

- ▶ Total expected reward for trajectory τ_i

$$R_{\tau_i} = \sum_{(s,u,s') \in \tau_i} \mathbb{E}_{r|s,u,s'}[r|s, u, s'] \quad (1)$$

- ▶ Maximum total expected reward to be $Q^*(s_i, u)$
- ▶ $Q(\cdot)$ is the action value function
- ▶ The recursive maximization equation, Bellman's equation,

$$Q^*(s_i, u) = \mathbb{E}_{s_j|s_i, u} [\mathbb{E}_{r|s_i, u, s_j} [r|s_i, u, s_j] + \gamma \max_{u' \in U(s_j)} Q_t(s_j, u')] \quad (2)$$



Q-Learning Background

- ▶ In many cases it is impossible to solve analytically, but can be formulated as iterative update

$$Q_{t+1}(s_i, u) = (1 - \alpha)Q_t(s_i, u) + \alpha[r_t + \gamma \max_{u' \in U(s_j)} Q_t(s_j, u')] \quad (4)$$

$$\lim_{t \rightarrow \infty} Q_t(s, u) = Q^*(s, u) \quad (5)$$

- ▶ The update equation has two parameters,
 - ▶ α is the Q-learning rate
 - ▶ γ is the discount factor



Q-Learning Background

- ▶ Q learning is *model free*
 - ▶ Agent can solve the task without explicitly constructing an estimate of the environmental dynamics
- ▶ Q learning is *off policy*
 - ▶ Agent can learn about optimal policies while exploring via non-optimal behavioral distribution
- ▶ Choose the behaviour distribution using ϵ -greedy strategy
 - ▶ Take a random action with probability ϵ
 - ▶ And greedy action $\max_{u \in U(s_t)} Q_t(s_t, u)$ with $1 - \epsilon$
- ▶ Experience replay
 - ▶ A memory of its explored paths and rewards
 - ▶ At a given interval agent samples from the memory and updates the Q -values



Designning Neural Network Architectures with Q-Learning

- ▶ The task of training a learning agent to sequentially choose neural network layers
- ▶ The agent will get a reward based on the validation accuracy
- ▶ Three main design choices
 - ▶ Reducing CNN layer definitions to simple state tuples
 - ▶ The set of layers the agent may pick
 - ▶ Balancing the size of the state-action space



The State Space

- ▶ Each state is a tuple of relevant layer parameters
 - ▶ Convolution (C)
 - ▶ Pooling (P)
 - ▶ Fully connected (FC)
 - ▶ Global average pooling (GAP)
 - ▶ Softmax (SM)
- ▶ Each layer has a parameter called layer depth and representation size (R-size)
- ▶ Convolutional nets progressively compress the representation of the original signal through pooling and convolution
- ▶ These layers in the state space may lead the agent on a trajectory where the intermediate signal representation gets reduced to a size that is too small for further processing.



The Action Space

- ▶ Restrict agent from taking certain actions.
 - ▶ Allow agent to terminate a path at any point
 - ▶ Allow transitions from state i to state with layer depth $i + 1$
 - ▶ Limit the number of fully connected (FC) layers max 2
 - ▶ A state s of type FC with number of neurons d may only transition to either a termination state or a state s' of type FC with number of neurons $d' \leq d$
 - ▶ An agent with state type of convolution (C) may transition to a state with any other layer type.
 - ▶ An agent with state type of pooling (P) may transition to a state with any other layer type other than P.



Q-Learning Training Procedure

- ▶ Learning rate $\alpha = 0.01$
- ▶ Discount factor $\gamma = 1$ not to over prioritize the short-term rewards
- ▶ Decrease the ϵ from 1 to 0.1 in steps
- ▶ Stop the agent $\epsilon = 0.1$ to obtain a stochastic final policy
- ▶ Main goal is to find well-performing model topologies can be ensembled to improve prediction performance
- ▶ During the training process maintain a replay dictionary storing
 - ▶ Network topology
 - ▶ Prediction performance on a validation set
- ▶ 100 models in the replay dictionary



Results




Figure: Q-Learning Performance on CNN accuracy for SVHN and CIFAR-10